

# Histograms

Nikos Hardavellas

Some slides/material from:

UToronto course by Andreas Moshovos

UIUC course by Wen-Mei Hwu and David Kirk

UCSB course by Andrea Di Blas

Universitat Jena by Waqar Saleem

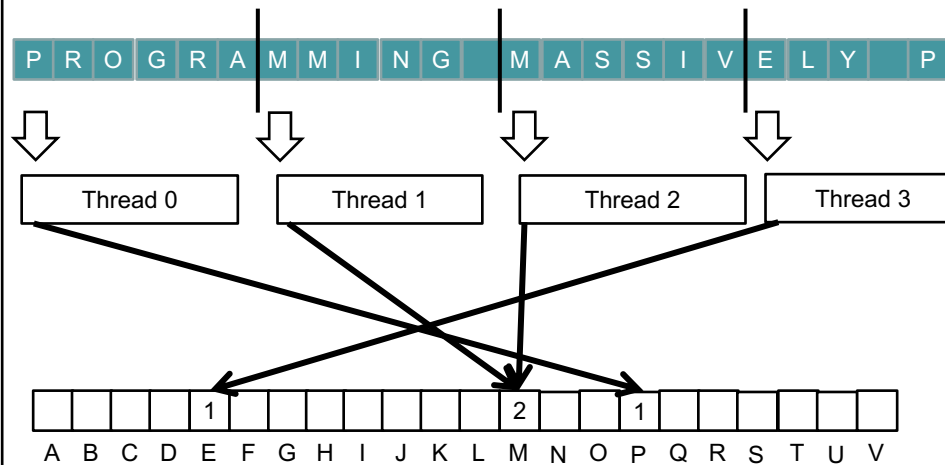
NVIDIA by Simon Green and many others

Patterson & Hennessy, Computer Organization: The HW/SW Interface  
Cedric Nugteren, Gert-Jan van den Braak, Henk Corporaal, Bart Mesman

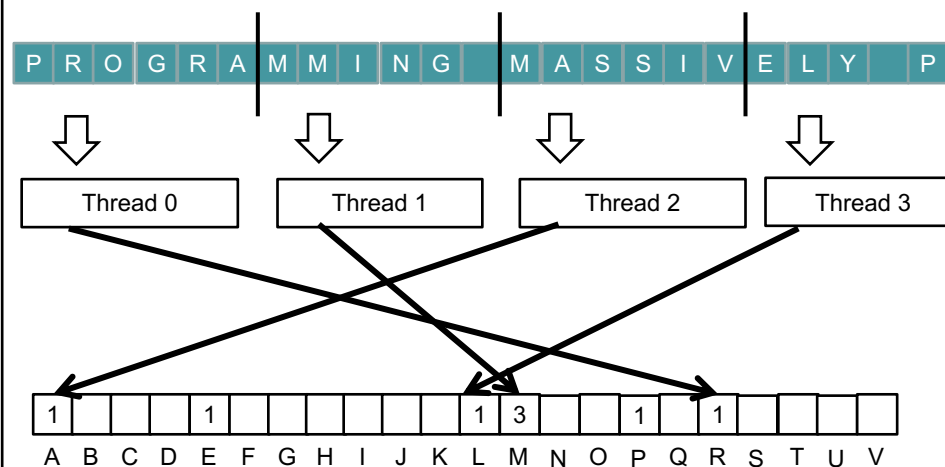
## A Histogram Example

- In the sentence  
“Programming Massively Parallel Processors”  
build a histogram of frequencies of each letter
- A(4), C(1), E(1), G(1), ...
- How do you do this in parallel?

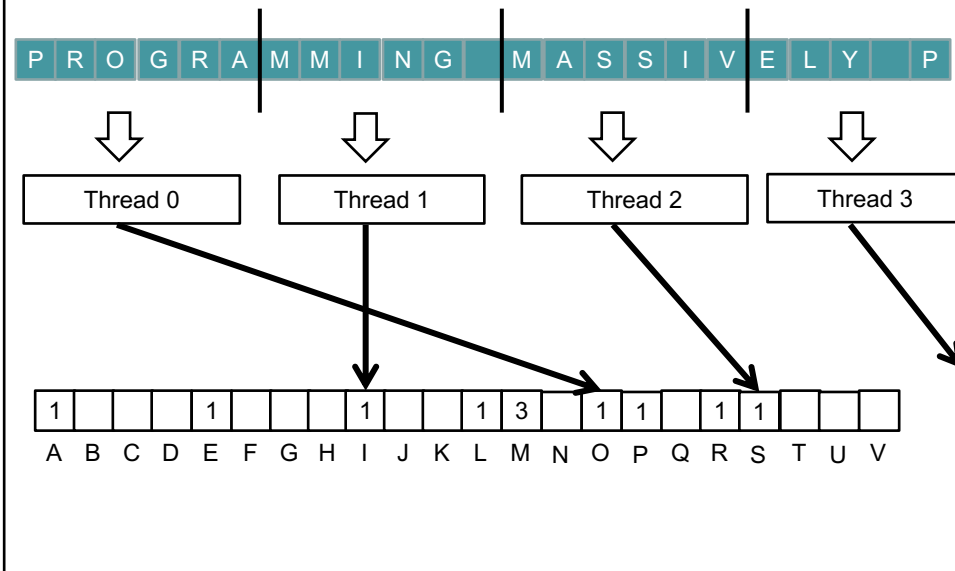
## Iteration #1



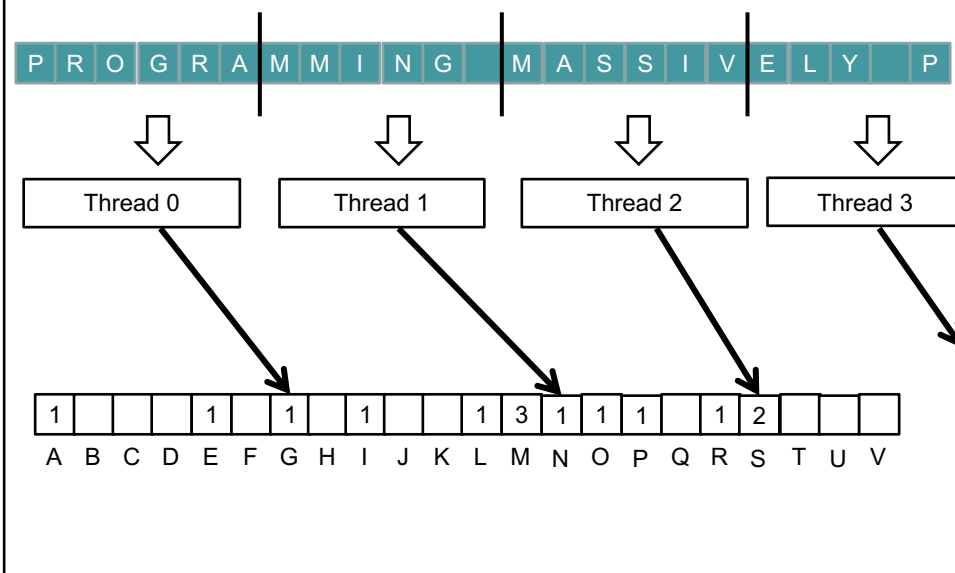
## Iteration #2



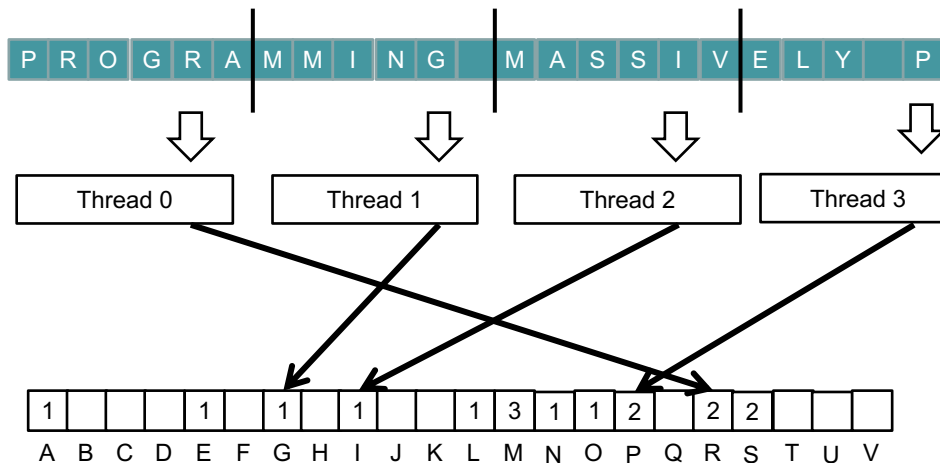
### Iteration #3



### Iteration #4



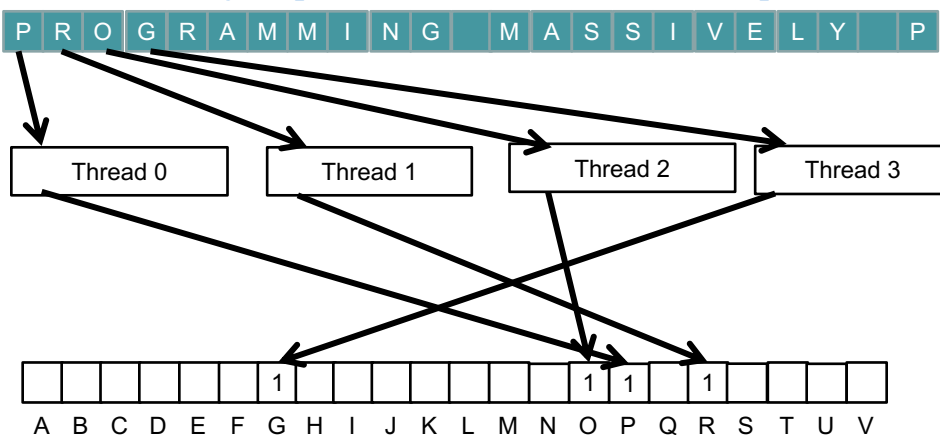
### Iteration #5



What is wrong with this algorithm?

### What is wrong with the algorithm?


- Reads from the input array are **not coalesced**
  - Assign inputs to each thread in a strided pattern



Races & conflicts in results array!!!

## Histogram

- E.g., Given an image calculate this:



- Distribution of values

## Sequential Algorithm

```
for (int i = 0; i < BIN_COUNT; i++)  
    result[i] = 0;
```

```
for (int i = 0; i < dataN; i++)  
    result[data[i]]++;
```

The challenge is that the write access pattern  
is data dependent

**No ordering of accesses to memory**

### Data Race

- |                    |          |
|--------------------|----------|
| • Thread 1         | Thread 2 |
| X++                | X++      |
| • X++ is really    |          |
| – tmp = Read x     |          |
| – tmp++            |          |
| – Write tmp into x |          |


### Data Race

- |                     |              |
|---------------------|--------------|
| • Start with X = 10 |              |
| • Thread 1          | Thread 2     |
| tmp = X (10)        | tmp = X (10) |
| tmp++ (11)          | tmp++ (11)   |
| X = tmp (11)        | X = tmp (11) |
| • Thread 1          | Thread 2     |
| tmp = X (10)        | ZZzzzzz      |
| tmp++ (11)          | ZZzzzzz      |
| X = tmp (11)        | ZZzzzzz      |
|                     | tmp = X (11) |
|                     | tmp++ (12)   |
|                     | X = tmp (12) |

## Parallel Strategy

- Distribute work across multiple blocks
  - Divide input data to blocks
- Each block process its own portion (**privatization**)
  - Multiple threads
    - $\# \text{pixels} / \# \text{threads}$  pixels per thread
  - Produces a partial histogram per thread
    - Could produce multiple histograms
    - One per thread  $\rightarrow$  no ordering problems (data races) here
- Merge all partial histograms
  - Produces the final histogram

## Parallel Strategy #1



- Input array in global memory
- Histogram array in shared memory
- One histogram per thread
- One column per possible pixel value

## Privatization

- Privatization is one of the most powerful and frequently used techniques for parallelizing applications
- Each thread works on its private data; merge later
- The operation needs to be
  - Commutative:  $x * y = y * x$
  - Associative:  $(x * y) * z = x * (y * z)$
  - Histogram add operation is associative and commutative

## What do we have control over?

```
result[ data[i] ] ++;
```

- **Data[ ]:**
  - *We control its access*
    - Can be accessed sequentially
  - Each element accessed only once
  - No races on data[ ] – it is read only
- **Result[ ]:**
  - Access is data-dependent
  - Each element may be accessed multiple times
  - Different pixels, read by different threads, same amplitude
- *We control memory placement*
  - Data[ ] in global memory (or constant?)
  - Result[ ] in shared memory
    - Needs to be small enough to fit in shared memory




## Sub-Histograms

- How big is each histogram (KB)?
  - Input value range: 0-255, 1 byte
  - Each histogram needs 256 entries
  - How many bytes per entry?
    - That's data dependent
  - Let's assume 32-bits or 4 bytes:  $256 * 4 = 1\text{KB}$  / histogram
- How many sub-histograms can we fit in shared mem?
  - Max shared mem in GTX680: 48KB
  - 48KB shared mem → only 1 warp (32KB in histograms)
  - Use only 1 warp per SM
  - Use only 32 out of 192 cores per SM
  - Unused hardware, low occupancy

## Sub-Histograms

- Let's try one histogram per block
  - Many threads per block
  - Ordering problem persists but within a block
  - However, threads within a block can **synchronize!!!**
  - **synchronization → no more race conditions**

## Parallel Strategy #2



- Input array in global memory
- Histogram array in shared memory
- One histogram per threadblock; **all threads in block update it**
- One column per possible pixel value
- **Each block updates global histogram at the end**


## Dataset Sketch for Global Histogram Update

Histogram in global memory:



But, it could be  $\#bins > \#threads/block$

→ Break histogram into  $\#threads/block$  pieces, work them iteratively



All threadblocks may be doing the update simultaneously!

## Algorithm Overview

- Step 1:
  - Initialize partial histogram
  - Each thread:
    - $s\_Hist[index] = 0$
    - $index += \text{threads per block}$ 
      - Until all bins are taken care of
- Step 2:
  - Generate partial histogram
  - Each thread:
    - read  $data[index]$
    - **update  $s\_Hist[]$**  ← conflicts possible (within block)
    - $index += \text{Total number of threads}$ 
      - Until all input/block is taken care of
- Step 3:
  - Update global histogram
  - Each thread
    - read  $s\_hist[index]$
    - **update global histogram** ← conflicts possible (across blocks)
    - $index += \text{threads per block}$ 
      - Until all bins are taken care of

## Simultaneous Updates?

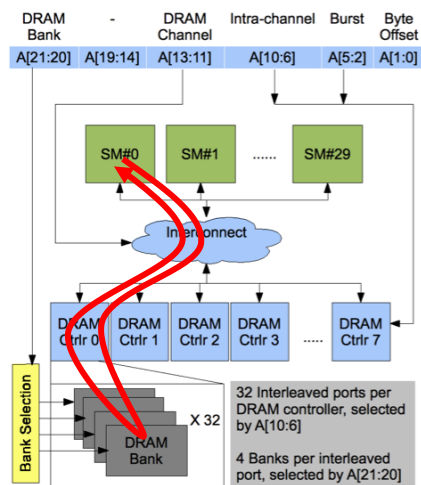
- Threads in a block:
  - update  $s\_Hist[]$
- All threads:
  - update global histogram
- Without special support this becomes:
  - register  $X = \text{value of } A$
  - $X ++$
  - $A = \text{register } X$
- This is a **read-modify-write** sequence

## The problem with simultaneous updates

- What if we do each step individually
 

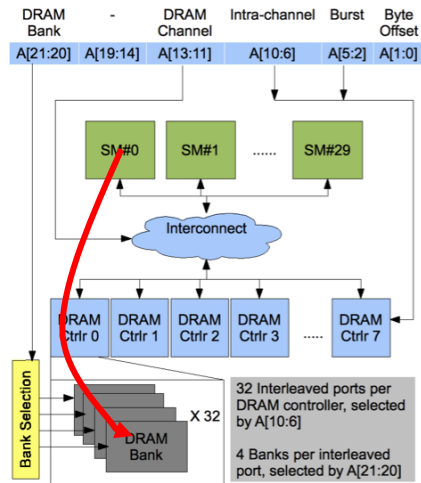
|                  |    |                 |    |
|------------------|----|-----------------|----|
| – r10 = mem[100] | 10 | r100 = mem[100] | 10 |
| – r10++          | 11 | r100++          | 11 |
| – mem[100] = r10 | 11 | mem[100] = r100 | 11 |
- But we really wanted 12
- What if we had 32 threads running in parallel?
- Starting with 10 we would want: 10+32
  - We may still get 11
- Need to think about this:
  - Special support: **Atomic operations**

## Atomic Operations on DRAM



- An atomic operation starts with a read, with a latency of a few hundred cycles

## Atomic Operations on DRAM



- An atomic operation starts with a read, with a latency of a few hundred cycles
- The atomic operation ends with a write, with a latency of a few hundred cycles
- During this whole time, no one else can access the location
- Similar for shared memory
- Latencies have improved!

## Atomic Operations

- Read-Modify-Write operations that are guaranteed to happen “atomically”
  - Produces the same result as if the sequence executed in isolation in time
  - Think of it as “serializing the execution” of all atomics
  - This is not what necessarily happens
    - This is how you should think about them

## Atomic Operations

- Supported both in Shared and Global memory
- Example:
  - `atomicAdd (pointer, value)`
  - does: `*pointer += value`
- Atomic Operations
  - Add, Sub, Inc, Dec
  - Exch, Min, Max, CAS (compare-and-swap)
  - Bitwise: And, Or, Xor
- Work with (unsigned) integers
- Exch works with floats as well

## `atomicExch`, `atomicMin`, `atomicMax`, `atomicCAS`

- `atomicExch (pointer, value)`
  - `tmp = * pointer`
  - `*pointer = value`
  - return `tmp`
- `atomicMin (pointer, value)` (max is similar)
  - `tmp = *pointer`
  - if (`*pointer > value`) `*pointer = value`
  - return `tmp`
- `atomicCAS (pointer, value1, value2)`
  - `tmp = *pointer`
  - if (`*pointer == value1`) `*pointer = value2`
  - return `tmp`

### atomicInc, atomicDec

- atomicInc (pointer, value)
  - tmp = \*pointer
  - if (\*pointer < value) (\*pointer)++
  - else \*pointer = 0
  - return tmp
- atomicDec (pointer, value)
  - tmp = \*pointer
  - if (\*pointer == 0 || \*pointer > value) \*pointer = value
  - else (\*pointer)--
  - return tmp
- Allow for wrap-around work queues

### atomicAnd, atomicOr, atomicXOR


- atomicAnd (pointer, value)
  - tmp = \*pointer
  - \*pointer = \*pointer & value
  - return tmp
- Others similar
- Now, you have two choices for the histogram:
  1. Use atomicInc(pointer, max\_value)
  2. Use atomicAdd(pointer, 1)

## Should you use Add or Inc ???

- `atomicAdd (pointer, value)`
  - `tmp = *pointer`
  - `*pointer = *pointer + value` Good!!!
  - `return tmp`
- `atomicInc (pointer, value)`
  - `tmp = *pointer`
  - `if (*pointer < value)`
  - `(*pointer)++` • Extra work ?
  - `else` • Thread Divergence ?
  - `*pointer = 0`
  - `return tmp`

## CUDA Implementation - Declarations

```
__global__ void histogram256Kernel  
(uint *d_Result, uint *d_Data, int dataN){  
  
    //Current global thread index  
    const int  
        globalTid = blockIdx.x * blockDim.x + threadIdx.x;  
  
    //Total number of threads in the compute grid  
    const int  
        numThreads = blockDim.x * gridDim.x;  
  
    __shared__ uint s_Hist[BIN_COUNT];
```



Why `uint *d_Data`?

Typically work with integers and floats  
These are 4-byte types  
Memory optimized for 4 bytes/access



### Clear partial histogram buffer

```
//Clear shared memory buffer for current block
before processing
    for ( int pos = threadIdx.x;
          pos < BIN_COUNT;
          pos += blockDim.x)
        s_Hist[pos] = 0;

    __syncthreads ();
    // All threads finished clearing out the
    // histogram
```

### Generate partial histogram

```
for (int pos = globalTid;
     pos < dataN;
     pos += numThreads) {

    uint data4 = d_Data[pos]; // coalesced
    // shared memory is word interleaved
    // read four pixels per thread with a load word

    atomicAdd (s_Hist + ((data4 >> 0) & 0xFFU),1);
    atomicAdd (s_Hist + ((data4 >> 8) & 0xFFU),1);
    atomicAdd (s_Hist + ((data4 >> 16) & 0xFFU),1);
    atomicAdd (s_Hist + ((data4 >> 24) & 0xFFU),1);
    // we are not using atomicInc which has a more
    // complex structure than atomicAdd
}

__syncthreads();
```

## Merge partial histogram with global histogram

```
for (int pos = threadIdx.x;
    pos < BIN_COUNT;
    pos += blockDim.x) {

    atomicAdd(d_Result + pos, s_Hist[pos]);
    // these operate on global memory
}
```

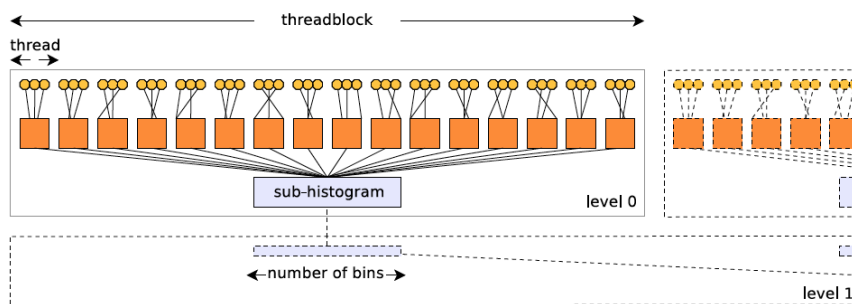
## Code overview

```
__global__ void histogram256Kernel (uint *d_Result, uint *d_Data,
int dataN){
    const int    globalTid = blockIdx.x * blockDim.x + threadIdx.x;
    const int    numThreads = blockDim.x * gridDim.x;
    __shared__ uint s_Hist[BIN_COUNT];
    for (int pos = threadIdx.x; pos < BIN_COUNT; pos += blockDim.x)
        s_Hist[pos] = 0;
    __syncthreads ();
    for (int pos = globalTid; pos < dataN; pos += numThreads){
        uint data4 = d_Data[pos]; // coalesced
        atomicAdd (s_Hist + (data4 >> 0) & 0xFFU, 1);
        atomicAdd (s_Hist + (data4 >> 8) & 0xFFU, 1);
        atomicAdd (s_Hist + (data4 >> 16) & 0xFFU, 1);
        atomicAdd (s_Hist + (data4 >> 24) & 0xFFU, 1);
    }
    __syncthreads ();
    for (int pos = threadIdx.x; pos < BIN_COUNT; pos += blockDim.x)
        atomicAdd(d_Result + pos, s_Hist[pos]);
}
```

## Discussion

- s\_Hist updates
  - Conflicts in shared memory
  - Data Dependent
  - 32-way conflicts possible and likely
- Is there an alternative?
  - One histogram per thread?
    - Not enough shared memory (if shared memory < 32 \* Hist\_Size)
    - Low occupancy (only 1 warp/block possible)
  - Load a tile of data in shared memory (or constant?)
    - Each thread produces a portion of the s\_Hist that maps onto the same bank
    - All threads read entire data tile


## Parallel Strategy #3



- Input array tile in shared memory
  - Every thread reads entire tile (in consecutive #threads \* 4 bytes chunks)
  - Wrap-around: in cycle  $i$ , where  $0 \leq i \leq 31$ , read  $\text{tile}[(i + \text{threadID}) \% 32]$
- Histogram array in shared memory
  - One column per possible pixel value
- One histogram per threadblock
  - Each thread responsible for a part of the sub-histogram
    - Only the bins that map to the same bank. Divergence?

## Map it to the same bank

Same color: sub-histogram bins for same thread



## This leads us to the Histogram Lab

- How many blocks?
- How many threads per block?
- How much shared memory used per thread?
  - One histogram per thread
    - ✓ Conflict/race free for partial histograms
    - ✗ But too few threads and blocks → Low occupancy
  - One histogram per threadblock
    - ✓ Higher occupancy
    - ✗ But conflicts in updates → atomics, bank conflicts
  - One histogram per threadblock, interleave among warp threads
    - ✓ Privatize portion of histogram per thread → no conflicts
    - ✗ But must read the input 32 times, (some) thread divergence
- Atomics or privatization? Conflicts or divergence? Low occupancy or conflicts? Read the input once or 32 times?
  - What is the right balance for histogram computation?
  - What other optimizations can you think of?